

Carla simulator implementation notes

Wang Xiang

October 13, 2022

Contents

1	Installation	1
2	Sensors configuration	2
2.1	Automatic data generation	2
2.2	Manual control in simulator	3
2.3	Sensor types in Carla	4
3	Road marking segmentation dataset preparation	4
3.1	Exporting front view camera images	4
3.2	Exporting BEV images	5
3.3	Post processing to build a segmentation dataset	5
3.3.1	Introduction to Ladybug5 camera	6
4	Public benchmarks about autonomous driving	6
4.1	Kitti	6
4.1.1	Road segmentation	6
4.1.2	Object detection	6
4.2	nuScenes	6
4.2.1	nuImages	7
4.2.2	nuReality	7
4.3	Waymo	7
4.4	ApolloScape	7
4.5	Soda-2d	7
5	Public datasets about medical segmentation	7

1 Installation

Official documentation is at: carla.readthedocs.io. Youtube official account is *CARLA Simulator*. The version that I installed is 0.9.13. The video on Youtube with title *CARLA 0.9.13* introduces new features of this version, including instance segmentation sensor that can generate instance segmentation ground truth mask. This sensor and semantic segmentation sensor may be useful in my project.

2 Sensors configuration

2.1 Automatic data generation

Youtube video: *An in depth look at CARLA's sensors*. Github repositories:

- github.com/Ozzyz/carla-data-export,
- github.com/jedeschaud/kitti-carla-simulator,
- github.com/jedeschaud/paris-carla-simulator,

SensorsCamera.py mainly referred the Youtube video mentioned above. Achievement: 1) I can run *SensorsCamera.py* successfully and see the car in a scene on the road. 2) I can also run *CameraSensorsType.py* successfully and see images from 6 different sensors. The vehicle uses autopilot and drives normally on the road. 3) I can also run *LidarRadarSensors3d.py* successfully and see RGB images, lidar and radar point clouds simultaneously.

Problem 1: over exposure of camera sensor. Attempt 1: change the weather parameters. The original weather parameters are

```
cloudiness=20.0, precipitation=0.0, precipitation_deposits=0.0,
wind_intensity=10.0, sun_azimuth_angle=300.0, sun_altitude_angle=45.0,
fog_density=2.0, fog_distance=0.75, fog_falloff=0.1, wetness=0.0,
scattering_intensity=1.0, mie_scattering_scale=0.03,
rayleigh_scattering_scale=0.0331
```

I turned all parameters besides *sun_azimuth_angle*, *sun_altitude_angle* to 0. When I turned *sun_altitude_angle* = -20.0, the view of RGB camera images appears normal sometimes when the buildings cover the sky, it corresponds to night scenes. But the view isn't normal when the buildings don't cover the sky, even if I set *sun_altitude_angle* = -60.0. Conclusion: the problem probably lies on the sky. Similar issues on Github are:

- github.com/carla-simulator/carla/issues/3417, .../4438, .../4527,

Another approach: change camera settings. Default camera settings are stored in *CamerargbAttributes.txt*. A useful command is

```
camera_bp.set_attribute('enable_postprocess_effects', 'False')
```

After that, both day scenes and night scenes look better. But the sky still has pure white color. Bad news is that all attributes of *sensor.camera.rgb* are the same on both Windows and Ubuntu systems.

Initially, each time when I launch *CarlaUE4.sh*, the following warning appears, and I ran the second line's command. It changes the file */proc/sys/dev/i915/perf_stream_paranoid* from 1 to 0. I don't know if it has any further unpredicted effects. The *sysctl* command displays and modifies kernel parameters in the directory */proc/sys*.

```
MESA-INTEL: warning: Performance support disabled, consider sysctl dev.
i915.perf_stream_paranoid=0
sudo sysctl dev.i915.perf_stream_paranoid=0
```

After comparing several fonts and styles of latex environment *lstlistings*, I found the following setting best meets my preference.

```
frame = single, basicstyle = \ttfamily\small, columns = fixed, breaklines = true
```

I also tried to install *Carla-0.9.13* on Windows system. Prerequisite: DirectX runtime. Running `dxwebsetup.exe` isn't successful, I installed it by downloading its full setup package. This time there's no bug on sky display. *Python 3.7* is required, note that its correct version is important, and I installed *3.7.9*. Run *manual_control.py* successfully on Windows system. Note that in order to import carla, I wrote two lines of code in *testcarlaiimport.py* and copy it every time that I need to import carla. Other requirements are listed as follows:

```
pip install pygame numpy opencv-python==4.2.0.34 open3d matplotlib
```

Good news is that I don't need to fix *Carla*'s bug of sky and RGB camera on Ubuntu system. I can use *Carla-0.9.13* on Windows system instead. An unpredicted advantage is that it runs faster on Windows.

2.2 Manual control in simulator

```
conda activate carla
cd PythonAPI/examples
python manual_control.py
```

Notable required packages in *requirements.txt* are *pygame* and *open3d*, neither of the two packages are installed in system python, conda base and openmmlab environments. I can run *manual_control.py* successfully and change sensors by pressing 1–9 keys. The sensors corresponding to these keys are listed as follows:

- 1: Camera RGB. Need to fix the bug of over exposure. Strange phenomenon is that the first second's scene after switching to RGB camera is slightly better.
- 2: Camera depth (raw). Displays normally. What is the relation between rendered color and depth?
- 3: Camera depth (grayscale). Displays normally.
- 4: Camera depth (logarithmic grayscale). Displays normally.
- 5: Camera semantic segmentation (raw).
- 6: Camera semantic segmentation (CityScapes Palette). Displays normally.
- 7: Camera instance segmentation (CityScapes Palette). Displays normally.
- 8: Camera instance segmentation (raw). Displays normally.
- 9: Lidar (ray-cast). Displays normally. Don't know what is the exact type of mounted lidar.

Some of other keys act normally but some keys don't act normally, and they are listed as follows:

- Backspace: change vehicle. Acts normally. Drawback is that the response time is approximately 7 seconds after pressing the key, and it's a little bit long.
- M: toggle manual transmission. I have to switch to manual transmission to drive the car, I don't know how to drive in automatic transmission mode.
- ,: gear down; .: gear up. I can switch to gear 1 – 6 or *R*. Pressing *W* is effective only when the gear is not *N*. It is weird that the gear number has no upper limit.

- W: throttle. When the gear is *R*, press *W* to drive backwards.
- O: open/close all doors of vehicle. Acts normally.
- R: toggle recording images to disk. Acts normally. I tested instance segmentation camera and RGB camera, the output images are .png files in subfolder `_out/`. Instance segmentation images have good quality, but RGB camera images still have over exposure. Frequency of exporting images is unknown, but the images are named using 8-digits integers, and I guess it's near 3Hz.

I can run `manual_control.py` in any directory besides `PythonAPI/examples`. In order to fix the over exposure bug, I decide to change the RGB camera configuration code in a copy of `manual_control.py` referring to github.com/carla-simulator/carla/issues/3634, but it turns out to be a failed attempt. Why?

Good news: while running `manual_control.py`, all keys act normally in Windows system.

2.3 Sensor types in Carla

Following its official account on Youtube, I wrote `CameraSensorsType.py` to explore different kinds of sensors in Carla. They are:

```
sensor.other.gnss, sensor.lidar.ray_cast, sensor.camera.semantic_segmentation,
sensor.other.radar, sensor.other.lane_invasion, sensor.camera.instance_segmentation,
sensor.camera.rgb, sensor.other.rss, sensor.camera.optical_flow,
sensor.lidar.ray_cast_semantic, sensor.other.obstacle, sensor.other.imu,
sensor.camera.depth, sensor.camera.dvs, sensor.other.collison,
```

Comments: 1) It has built in lidar semantic sensor, but I don't know how to use it currently.

2) DVS stands for *dynamic vision sensor*, also known as *event camera*, *neuromorphic camera* or *silicon retina*, is an imaging sensor that responds to local changes in brightness. Event cameras don't capture images using a shutter as conventional frame cameras do. Instead, each pixel inside an event camera operates independently and asynchronously, reporting changes in brightness as they occur, and staying silent otherwise.

3) What is an optical flow sensor? It is a vision sensor capable of measuring optical flow or visual motion and outputting a measurement based on optical flow. Optical flow is the pattern of apparent motion of objects, surfaces, and edges in a visual scene caused by the relative motion between an observer and a scene. It can also be defined as the distribution of apparent velocities of movement of brightness pattern in an image.

3 Road marking segmentation dataset preparation

3.1 Exporting front view camera images

Currently the only known example way of exporting images is to press *R* in `manual_control.py`, so I may refer to this code. The second known example is `tutorial.py`. I modified the callback functions of *rgb*, *semantic*, *instance*, *depth* camera sensors in `CameraSensorsType.py`, and successfully exported the above 4 sensors' images.

3.2 Exporting BEV images

One possible way is to rotate the spectator. Reference:

- *medium.com: From semantic segmentation to semantic bird's eye view in the Carla simulator, by Maciek Dziubiński;*
- *Youtube: Semantic bird's view, by Acta Schola Automata Polonica;*
- *NEAT: Neural attention fields for end-to-end autonomous driving, by Andreas Geiger, ICCV 2021; code repository: github.com/autonomousvision/neat;*
- *carla.readthedocs.io/en/0.9.13/python_api/#carlatransform, .../[#carla.Rotation](https://carla.readthedocs.io/en/0.9.13/python_api/#carla.Rotation);*

I'm mainly interested in the way how BEV semantic segmentation ground truth is generated. If I can walk through this pipeline, I believe that generating BEV rgb camera images can be done. Key points are two classes: *carla.Transform* and *carla.Rotation*. *carla.Transform* consists of two instance variables: location (*carla.Location* class) and rotation (*carla.Rotation* class). Notice that *Unreal Engine*'s coordinate system adopted left-handed convention. Correct configuration for BEV camera is *carla.Rotation(pitch=-90)*, unit of Euler angles is degree. Now I can modify the poses of both spectator and camera.

Problem 1 (Reinstall ros-kinetic-desktop-full for Lianzhong).

I backtraced dependency relations through the following commands:

```
sudo apt install ros-kinetic-desktop-full
sudo apt install ros-kinetic-desktop
sudo apt install ros-kinetic-viz
sudo apt install ros-kinetic-rqt-common-plugins
sudo apt install ros-kinetic-rqt-image-view
sudo apt install ros-kinetic-opencv3
sudo apt install libvtk6-qt-dev
sudo apt install libvtk6-dev
sudo apt install libgdal-dev
```

The message of the last two dependencies, and solution to fix this bug are as follows:

```
sudo apt install libgeos-dev
sudo aptitude install ros-kinetic-desktop-full
The following packages have unmet dependencies:
  libgeos-dev : Depends: libgeos-c1v5 (= 3.5.0-1ubuntu2) but 3.5.1-3~
                  xenial0 is to be installed
sudo apt install libgeos-c1v5=3.5.0-1ubuntu2
```

Now I know that *aptitude* command works better than *apt* sometimes. In this case, it reports the key dependent package with unmet version through backtracing automatically. Main reference: <https://blog.csdn.net/HelloJinYe/article/details/109104249>.

3.3 Post processing to build a segmentation dataset

Several related questions are listed as follows:

- *How to change the pose of lidar mounted on vehicle? It should be pitch=15.*
- *How to export BEV images of lidar intensity?*

- *How to hide ego vehicle during data acquisition?*
- *How to plan the path that the vehicle drives through?*
- *How to configure the size of exported images?*
- *How to configure the city that the vehicle is driving in?*

According to github.com/carla-simulator/carla/issues/4261, currently it is not possible to hide ego vehicle in python API. I should modify the source of the simulator if I am willing to do so.

Update 2022-09-30: I can run the code of *paris-carla-simulator* successfully on Windows system and export generated images now. Reference: [1]. The original source code exports raw semantic segmentation images, I converted them using CityScapes palette. The sizes of both RGB camera images and semantic camera images are 2048×2048 . RGB cameras and semantic cameras are mounted on 6 different poses. The paper says their mobile system equips *Velodyne HDL32* lidar and *Ladybug5* camera. Next, I want to understand its source code. What can I learn from its source code? *recording.log* is too large and I don't know how to use it, so I deleted it.

3.3.1 Introduction to Ladybug5 camera

4 Public benchmarks about autonomous driving

4.1 Kitti

4.1.1 Road segmentation

4.1.2 Object detection

4.2 nuScenes

The *nuScenes* dataset is a large-scale autonomous driving dataset with *3d object annotations*. It features:

- Full sensor suite (1 lidar, 5 radars, 6 cameras, imu, gps);
- 1000 scene of 20s each, 1.4m camera images, 390k lidar sweeps;
- Two cities: Boston and Singapore, left versus right hand traffic, detailed map information;
- 1.4m 3d bounding boxes manually annotated for 23 object classes;
- Attributes such as visibility, activity and pose;
- 1.1b lidar points manually annotated for 32 classes;

Problem 2 (Detection). Average translation error (ATE): Euclidean center distance in 2D in meters;

Average scale error (ASE): calculated as 1-IOU after aligning centers and orientation;

Average orientation error (AOE): Smallest yaw angle difference between prediction and ground truth in radians. Orientation error is evaluated at 360 degree for all classes except barriers where it is only evaluated at 180 degrees. Orientation errors for cones are ignored;

Average velocity error (AVE): Absolute velocity error in m/s. Velocity error for barriers and cones are ignored;

Average attribute error (AAE): Calculated as $1 - \text{acc}$, where acc is the attribute classification accuracy. Attribute error for barriers and cones are ignored;

The TP metrics are defined per class, and we then take a mean over classes to calculate mATE, mASE, mAOE, mAVE and mAAE.

nuScenes detection score (NDS): We consolidate the above metrics by computing a weighted sum: mAP, mATE, mASE, mAOE, mAVE and mAAE. As a first step we convert the TP errors to TP scores as $TP_{score} = \max(1 - TP_{error}, 0.)$. We then assign a weight of 5 to mAP and 1 to each of the 5 TP scores and calculate the normalized sum.

Lidar track: Only lidar input allowed. Vision track: Only camera input allowed. Open track: Any sensor input allowed.

4.2.1 nuImages

nuImages is a large-scale autonomous driving dataset with *image-level 2d annotations*. It features:

- 93k video clips of 6s each (150h of driving);
- 93k annotated and 1.1m un-annotated images;
- Two cities: Boston and Singapore;
- The same proven sensor suite as in nuScenes;
- Images mined for diversity;
- 800k annotated foreground objects with 2d bounding boxes and instance masks;
- 100k 2d semantic segmentation masks for background classes;

4.2.2 nuReality

The *nuReality* environment is a novel virtual reality environment designed to test the interactions between pedestrians and autonomous vehicles.

4.3 Waymo

4.4 Apolloscape

4.5 Soda-2d

5 Public datasets about medical segmentation

celltrackingchallenge.net/2d-datasets. The first processed dataset is simulated nuclei dataset. Wrote program to analyse the histogram of pixel values. Folder *01/* contains 65 images, folder *02/* contains 150 images, their histograms are written in *pixelhistogram01.txt* and *pixelhistogram02.txt* respectively. Raw images are normalized by the following steps:

```
lowerbound = 75; newupperbound = 127
imagepng1 = np.maximum(imagepng, lowerbound) - lowerbound; imagepng1 = np.
    minimum(imagepng1, newupperbound).astype(np.uint8)
```

Notice that in the above steps the relative pixel intensity is not changed. I'm focusing on the segmentation task. The next question is, how is its ground truth labelled? Use custom palette to visualize the ground truth mask. Relevant functions are `cv.LUT()`, `cv.applyColorMap()`, notice that I manually changed `lut[0] = 0` so that the background color is black. Ground truth images provide instance segmentation mask, in addition, same cell in different frames are labelled with the same integer, background label is 0. But I still don't know what information does *man_track.txt* give.

Now I think data preparation is done. Maybe I can start doing semantic and instance segmentation now. Is 65+150 pairs of images and labels enough for training a model? Challenge dataset contains 110+138 images without annotation. Real datasets' annotations are classified as gold truth and silver truth. The second processed dataset is DicHela dataset. Major difference is the following thresholds:

```
lowerbound = 50; newupperbound = 127
```

6 Labeling tools

ApolloScape dataset recommended *label-studio*. At first I tried to install it using system's pip, but uninstalled those packages soon. I learned that `gedit` can search regular expression in a file, I used the following regular expression to match the version identifiers after the packages' names.

```
-[0123456789]*[.][0123456789]*[.][0123456789]
```

References

- [1] Jean-Emmanuel Deschaud, Paris-Carla-3D: A real and synthetic outdoor point cloud dataset for challenging tasks in 3D mapping, Remote Sensing, 2021.
code repository: github.com/jedeschaud/paris_carla_simulator.